

# Block 03 - Data Handling

*Marcelo*

*20 November 2017*

```
#rm(list = ls())
library(dplyr)
library(nycflights13)
```

dplyr aims to provide a function for each basic verb of data manipulating:

- filter() (and slice())
- arrange()
- select() (and rename())
- distinct()
- mutate() (and transmute())
- summarise()
- sample\_n() and sample\_frac()

```
class(flights)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
dfFlights <- as.data.frame(flights)
```

```
flights[flights$month == 7 &
        flights$day == 1, ]
```

```
## # A tibble: 966 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     7     1         1           2029         212     236
## 2  2013     7     1         2           2359          3     344
## 3  2013     7     1        29           2245        104     151
## 4  2013     7     1        43           2130        193     322
## 5  2013     7     1        44           2150        174     300
## 6  2013     7     1        46           2051        235     304
## 7  2013     7     1        48           2001        287     308
## 8  2013     7     1        58           2155        183     335
## 9  2013     7     1       100           2146        194     327
## 10 2013     7     1       100           2245        135     337
## # ... with 956 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
system.time(filter(flights, month == 7, day == 1))
```

```
##   user  system elapsed
##   0.01   0.00   0.01
```

```
system.time(filter(dfFlights, dep_delay < 0 & arr_delay > 0))
```

```
##   user  system elapsed
##   0.02   0.00   0.02
```

```
#arrange(flights, year, month, day)
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641             900          1301    1242
## 2  2013     6    15    1432            1935          1137    1607
## 3  2013     1    10    1121            1635          1126    1239
## 4  2013     9    20    1139            1845          1014    1457
## 5  2013     7    22     845            1600          1005    1044
## 6  2013     4    10    1100            1900           960    1342
## 7  2013     3    17    2321             810           911     135
## 8  2013     7    22    2257             759           898     121
## 9  2013    12     5     756            1700           896    1058
## 10 2013     5     3    1133            2055           878    1250
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

all the following outputs are identical

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
select(flights, 1:3)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

## Exercices: Subset, Filter and Arrange data frames

### 1 and 2

```
dataPath <- "./Datasets/Schweiz/"
filesToLoad <- list.files(path = dataPath, full.names = TRUE)
filesToLoad

## [1] "./Datasets/Schweiz/Schweiz1.txt" "./Datasets/Schweiz/Schweiz2.txt"
## [3] "./Datasets/Schweiz/Schweiz3.txt" "./Datasets/Schweiz/Schweiz4.txt"
```

### 3

```
dat1 <- read.csv2(file = filesToLoad[1])
str(dat1, vec.len = 1)

## 'data.frame':   47 obs. of  7 variables:
##  $ Province      : Factor w/ 47 levels "Aigle","Aubonne",...: 8 9 ...
##  $ Fertility      : num  80.2 83.1 ...
##  $ Agriculture    : num  17 45.1 ...
##  $ Examination    : int   15 6 ...
##  $ Education      : int   12 9 ...
##  $ Catholic       : num   9.96 ...
##  $ Infant.Mortality: num  22.2 22.2 ...

dat2 <- read.csv(file = filesToLoad[2])

dat1 <- read.table(file= filesToLoad[1],
                  header = TRUE,
                  sep = ";",
                  dec = ",",
                  stringsAsFactors = T)
head(dat1)

##      Province Fertility Agriculture Examination Education Catholic
## 1  Courtelary    80.2      17.0          15         12      9.96
```

```
## 2      Delemont      83.1      45.1          6          9      84.84
## 3 Franches-Mnt      92.5      39.7          5          5      93.40
## 4      Moutier      85.8      36.5         12          7      33.77
## 5      Neuveville    76.9      43.5         17         15       5.16
## 6      Porrentruy    76.1      35.3          9          7      90.57
## Infant.Mortality
## 1          22.2
## 2          22.2
## 3          20.2
## 4          20.3
## 5          20.6
## 6          26.6
```

```
dat2 <- read.table(file= filesToLoad[2],
                    header = TRUE,
                    sep = ",",
                    dec = ".",
                    stringsAsFactors = T)
head(dat2)
```

```
##      Province Fertility Agriculture Examination Education Catholic
## 1  Courtelary    80.2      17.0          15          12     9.96
## 2   Delemont    83.1      45.1           6           9    84.84
## 3 Franches-Mnt  92.5      39.7           5           5    93.40
## 4   Moutier    85.8      36.5          12           7    33.77
## 5  Neuveville   76.9      43.5          17          15     5.16
## 6  Porrentruy   76.1      35.3           9           7    90.57
## Infant.Mortality
## 1          22.2
## 2          22.2
## 3          20.2
## 4          20.3
## 5          20.6
## 6          26.6
```

```
dat3 <- read.table(file= filesToLoad[3],
                    header = TRUE,
                    sep = "_",
                    dec = "-",
                    stringsAsFactors = T)
head(dat3)
```

```
##      Province Fertility Agriculture Examination Education Catholic
## 1  Courtelary    80.2      17.0          15          12     9.96
## 2   Delemont    83.1      45.1           6           9    84.84
## 3 Franches-Mnt  92.5      39.7           5           5    93.40
## 4   Moutier    85.8      36.5          12           7    33.77
## 5  Neuveville   76.9      43.5          17          15     5.16
## 6  Porrentruy   76.1      35.3           9           7    90.57
## Infant.Mortality
## 1          22.2
## 2          22.2
## 3          20.2
## 4          20.3
## 5          20.6
## 6          26.6
```

```
dat4 <- read.table(file= filesToLoad[4],
                  header = TRUE,
                  sep = "\t",
                  dec = ",",
                  stringsAsFactors = T)
head(dat4)
```

```
##      Province Fertility Agriculture Examination Education Catholic
## 1  Courtelary      80.2        17.0           15         12      9.96
## 2   Delemont      83.1        45.1            6          9     84.84
## 3 Franches-Mnt    92.5        39.7            5          5     93.40
## 4    Moutier      85.8        36.5           12          7     33.77
## 5  Neuveville     76.9        43.5           17         15      5.16
## 6  Porrentruy     76.1        35.3            9          7     90.57
## Infant.Mortality
## 1              22.2
## 2              22.2
## 3              20.2
## 4              20.3
## 5              20.6
## 6              26.6
```

```
all(identical(dat1, dat2),
    identical(dat3, dat4),
    identical(dat1, dat3))
```

```
## [1] TRUE
```

```
datSub1 <- dat1[-(1:3), 1:6]
datSub2 <- dat1[-(4:10), c(1, 7)]
```

## 9.

Use the \* join functions from the dplyr package to reunite the two subsets to one data frame named datJoin. Read the help file for inner join. Which column should be used for joining? Your result should look like:

```
datJoin <- full_join(x = datSub1,
                    y = datSub2,
                    by = "Province")
str(datJoin, vec.len = 1)
```

```
## 'data.frame':  47 obs. of  7 variables:
## $ Province      : Factor w/ 47 levels "Aigle","Aubonne",...: 26 28 ...
## $ Fertility      : num  85.8 76.9 ...
## $ Agriculture    : num  36.5 43.5 ...
## $ Examination    : int   12 17 ...
## $ Education      : int    7 15 ...
## $ Catholic       : num  33.8 ...
## $ Infant.Mortality: num  NA NA ...
```

```
identical(dim(dat1), dim(datJoin))
```

```
## [1] TRUE
```

```
identical(dat1, datJoin)
```

```
## [1] FALSE
```

## Subset, Filter and Arrange data frames

```
data(swiss)
?swiss
```

```
## starting httpd help server ... done
```

Create a new column `swiss$Province` in which you store the name of the provinces corresponding to the observations. Delete the rownames afterwards.

```
swiss <- swiss %>%
  mutate(Province = rownames(swiss)) %>%
  select(7, 1:6)
swiss
```

##	Province	Fertility	Agriculture	Examination	Education	Catholic
## 1	Courtelay	80.2	17.0	15	12	9.96
## 2	Delemont	83.1	45.1	6	9	84.84
## 3	Franches-Mnt	92.5	39.7	5	5	93.40
## 4	Moutier	85.8	36.5	12	7	33.77
## 5	Neuveville	76.9	43.5	17	15	5.16
## 6	Porrentruy	76.1	35.3	9	7	90.57
## 7	Broye	83.8	70.2	16	7	92.85
## 8	Glane	92.4	67.8	14	8	97.16
## 9	Gruyere	82.4	53.3	12	7	97.67
## 10	Sarine	82.9	45.2	16	13	91.38
## 11	Veveyse	87.1	64.5	14	6	98.61
## 12	Aigle	64.1	62.0	21	12	8.52
## 13	Aubonne	66.9	67.5	14	7	2.27
## 14	Avenches	68.9	60.7	19	12	4.43
## 15	Cossonay	61.7	69.3	22	5	2.82
## 16	Echallens	68.3	72.6	18	2	24.20
## 17	Grandson	71.7	34.0	17	8	3.30
## 18	Lausanne	55.7	19.4	26	28	12.11
## 19	La Vallee	54.3	15.2	31	20	2.15
## 20	Lavaux	65.1	73.0	19	9	2.84
## 21	Morges	65.5	59.8	22	10	5.23
## 22	Moudon	65.0	55.1	14	3	4.52
## 23	Nyone	56.6	50.9	22	12	15.14
## 24	Orbe	57.4	54.1	20	6	4.20
## 25	Oron	72.5	71.2	12	1	2.40
## 26	Payerne	74.2	58.1	14	8	5.23
## 27	Paysd'enhaut	72.0	63.5	6	3	2.56
## 28	Rolle	60.5	60.8	16	10	7.72
## 29	Vevey	58.3	26.8	25	19	18.46
## 30	Yverdon	65.4	49.5	15	8	6.10
## 31	Conthey	75.5	85.9	3	2	99.71
## 32	Entremont	69.3	84.9	7	6	99.68
## 33	Herens	77.3	89.7	5	2	100.00
## 34	Martigwy	70.5	78.2	12	6	98.96
## 35	Monthey	79.4	64.9	7	3	98.22

## 36	St Maurice	65.0	75.9	9	9	99.06
## 37	Sierre	92.2	84.6	3	3	99.46
## 38	Sion	79.3	63.1	13	13	96.83
## 39	Boudry	70.4	38.4	26	12	5.62
## 40	La Chauxdfnd	65.7	7.7	29	11	13.79
## 41	Le Locle	72.7	16.7	22	13	11.22
## 42	Neuchatel	64.4	17.6	35	32	16.92
## 43	Val de Ruz	77.6	37.6	15	7	4.97
## 44	ValdeTravers	67.6	18.7	25	7	8.65
## 45	V. De Geneve	35.0	1.2	37	53	42.34
## 46	Rive Droite	44.7	46.6	16	29	50.43
## 47	Rive Gauche	42.8	27.7	22	29	58.33
##	Infant.Mortality					
## 1		22.2				
## 2		22.2				
## 3		20.2				
## 4		20.3				
## 5		20.6				
## 6		26.6				
## 7		23.6				
## 8		24.9				
## 9		21.0				
## 10		24.4				
## 11		24.5				
## 12		16.5				
## 13		19.1				
## 14		22.7				
## 15		18.7				
## 16		21.2				
## 17		20.0				
## 18		20.2				
## 19		10.8				
## 20		20.0				
## 21		18.0				
## 22		22.4				
## 23		16.7				
## 24		15.3				
## 25		21.0				
## 26		23.8				
## 27		18.0				
## 28		16.3				
## 29		20.9				
## 30		22.5				
## 31		15.1				
## 32		19.8				
## 33		18.3				
## 34		19.4				
## 35		20.2				
## 36		17.8				
## 37		16.3				
## 38		18.1				
## 39		20.3				
## 40		20.5				
## 41		18.9				

```
## 42          23.0
## 43          20.0
## 44          19.5
## 45          18.0
## 46          18.2
## 47          19.3
```

Identify the five provinces with the lowest percentage of males involved in agriculture as occupation.

```
swiss %>%
  arrange(Agriculture) %>%
  select(Province, Agriculture) %>%
  slice(1:5)
```

```
## # A tibble: 5 x 2
##   Province Agriculture
##   <chr>         <dbl>
## 1 V. De Geneve      1.2
## 2 La Chauxdfnd      7.7
## 3 La Vallee        15.2
## 4 Le Locle         16.7
## 5 Courtelary       17.0
```

Find the five provinces with the highest percentage of catholic population.

```
swiss %>%
  arrange(desc(Catholic)) %>%
  select(Province, Catholic) %>%
  slice(1:5)
```

```
## # A tibble: 5 x 2
##   Province Catholic
##   <chr>         <dbl>
## 1 Herens      100.00
## 2 Conthey     99.71
## 3 Entremont   99.68
## 4 Sierre      99.46
## 5 St Maurice  99.06
```

## Data Manipulation

```
rm(list = ls())
library(dplyr)
```

```
urlRedWine <-
"http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv"
urlWhiteWine <-
"http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"
readLines(urlRedWine, n = 5)
```



```
## [1] "\"fixed acidity\";\"volatile acidity\";\"citric acid\";\"residual sugar\";\"chlorides\";\"free s
## [2] "7.4;0.7;0;1.9;0.076;11;34;0.9978;3.51;0.56;9.4;5"
## [3] "7.8;0.88;0;2.6;0.098;25;67;0.9968;3.2;0.68;9.8;5"
## [4] "7.8;0.76;0.04;2.3;0.092;15;54;0.997;3.26;0.65;9.8;5"
## [5] "11.2;0.28;0.56;1.9;0.075;17;60;0.998;3.16;0.58;9.8;6"
```

```
readLines(urlWhiteWine, n = 5)
```

```
## [1] "\"fixed acidity\";\"volatile acidity\";\"citric acid\";\"residual sugar\";\"chlorides\";\"free s
## [2] "7;0.27;0.36;20.7;0.045;45;170;1.001;3;0.45;8.8;6"
## [3] "6.3;0.3;0.34;1.6;0.049;14;132;0.994;3.3;0.49;9.5;6"
## [4] "8.1;0.28;0.4;6.9;0.05;30;97;0.9951;3.26;0.44;10.1;6"
## [5] "7.2;0.23;0.32;8.5;0.058;47;186;0.9956;3.19;0.4;9.9;6"
```

```
redWine <- read.table(file = urlRedWine,
                     header = TRUE,
                     sep = ";",
                     dec = ".")
```

```
whiteWine <- read.table(file = urlWhiteWine,
                       header = TRUE,
                       sep = ";",
                       dec = ".")
```

```
redWine <- mutate(.data = redWine, color = "red")
whiteWine <- mutate(.data = whiteWine, color = "white")
```

```
wine <- bind_rows(redWine, whiteWine)
str(wine, vec.len = 2)
```

```
## 'data.frame':   6497 obs. of  13 variables:
## $ fixed.acidity      : num  7.4 7.8 7.8 11.2 7.4 ...
## $ volatile.acidity   : num  0.7 0.88 0.76 0.28 0.7 ...
## $ citric.acid        : num  0 0 0.04 0.56 0 ...
## $ residual.sugar     : num  1.9 2.6 2.3 1.9 1.9 ...
## $ chlorides          : num  0.076 0.098 0.092 0.075 0.076 ...
## $ free.sulfur.dioxide : num  11 25 15 17 11 ...
## $ total.sulfur.dioxide: num  34 67 54 60 34 ...
## $ density            : num  0.998 0.997 ...
## $ pH                 : num  3.51 3.2 3.26 3.16 3.51 ...
## $ sulphates          : num  0.56 0.68 0.65 0.58 0.56 ...
## $ alcohol            : num  9.4 9.8 9.8 9.8 9.4 ...
## $ quality            : int  5 5 5 6 5 ...
## $ color              : chr  "red" "red" ...
```

```
rm(redWine, whiteWine)
```

```
# apply condition
wine <- wine %>%
  mutate(alcoholClass = ifelse(test = alcohol < 10,
                              yes = "low",
                              no = ifelse(test = alcohol < 12,
                                          yes = "medium",
                                          no = "high"))))
```

```
# apply ordered factor
wine$alcoholClass <- factor(x = wine$alcoholClass,
```

```

        levels = c("low", "medium", "high"),
        ordered = TRUE)

#display aggregated table
table(wine$alcoholClass, wine$color)

##
##           red white
## low      680 1923
## medium   757 2162
## high     162  813

prop.table(table(wine$alcoholClass, wine$color), margin = 2)

##
##           red      white
## low    0.4252658 0.3926092
## medium 0.4734209 0.4414047
## high   0.1013133 0.1659861

```

6. Figure out how the same results can be obtained using dplyr functions:

```

wine %>%
  select(color, alcoholClass) %>%
  count(color, alcoholClass)

## # A tibble: 6 x 3
##   color alcoholClass     n
##   <chr>         <ord> <int>
## 1 red          low    680
## 2 red          medium 757
## 3 red          high   162
## 4 white        low   1923
## 5 white        medium 2162
## 6 white        high   813

# not quite right
wine %>%
  select(color, alcoholClass) %>%
  count(color, alcoholClass) %>%
  group_by(color) %>% mutate(ratio = n / sum(n))

## # A tibble: 6 x 4
## # Groups:   color [2]
##   color alcoholClass     n    ratio
##   <chr>         <ord> <int>   <dbl>
## 1 red          low    680 0.4252658
## 2 red          medium 757 0.4734209
## 3 red          high   162 0.1013133
## 4 white        low   1923 0.3926092
## 5 white        medium 2162 0.4414047
## 6 white        high   813 0.1659861

```

7. Calculate the mean values for alcohol, density and pH for each combination of quality and color using the aggregate function.

```
wineAgg <- wine %>%
  select(quality, color, alcohol, density, pH)

aggregate(x = wineAgg[, -(1:2)],
          by = list(wineAgg$quality, wineAgg$color),
          FUN = mean)
```

##	Group.1	Group.2	alcohol	density	pH
## 1	3	red	9.955000	0.9974640	3.398000
## 2	4	red	10.265094	0.9965425	3.381509
## 3	5	red	9.899706	0.9971036	3.304949
## 4	6	red	10.629519	0.9966151	3.318072
## 5	7	red	11.465913	0.9961043	3.290754
## 6	8	red	12.094444	0.9952122	3.267222
## 7	3	white	10.345000	0.9948840	3.187500
## 8	4	white	10.152454	0.9942767	3.182883
## 9	5	white	9.808840	0.9952626	3.168833
## 10	6	white	10.575372	0.9939613	3.188599
## 11	7	white	11.367936	0.9924524	3.213898
## 12	8	white	11.636000	0.9922359	3.218686
## 13	9	white	12.180000	0.9914600	3.308000

## Block 03.02 - 27.11.17

```
rm(list = ls())
```

```
require(nycflights13)
require(dplyr)
```

Rename at select

```
select(flights, depTime = dep_time)
```

```
## # A tibble: 336,776 x 1
##   depTime
##   <int>
## 1     517
## 2     533
## 3     542
## 4     544
## 5     554
## 6     554
## 7     555
## 8     557
## 9     557
## 10    558
## # ... with 336,766 more rows

rename (returns full DataFrame)
```

```
rename(flights, depTime = dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month day depTime sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1   1     517           515         2     830
## 2  2013     1   1     533           529         4     850
## 3  2013     1   1     542           540         2     923
## 4  2013     1   1     544           545        -1    1004
## 5  2013     1   1     554           600        -6     812
## 6  2013     1   1     554           558        -4     740
## 7  2013     1   1     555           600        -5     913
## 8  2013     1   1     557           600        -3     709
## 9  2013     1   1     557           600        -3     838
## 10 2013     1   1     558           600        -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

analogous to unique, but for whole dataFrame

```
distinct(flights, tailnum)
```

```
## # A tibble: 4,044 x 1
##   tailnum
##   <chr>
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## # ... with 4,034 more rows
```

list all flights connections (unique origin AND destination)

```
distinct(select(flights, origin, dest))
```

```
## # A tibble: 224 x 2
##   origin dest
##   <chr> <chr>
## 1 EWR   IAH
## 2 LGA   IAH
## 3 JFK   MIA
## 4 JFK   BQN
## 5 LGA   ATL
## 6 EWR   ORD
## 7 EWR   FLL
## 8 LGA   IAD
## 9 JFK   MCO
## 10 LGA   ORD
```

```
## # ... with 214 more rows
```

```
mutate
```

```
mutate(flights,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60) %>%  
  select(gain, speed, distance, air_time, arr_delay, dep_delay)
```

```
## # A tibble: 336,776 x 6
```

```
##   gain    speed distance air_time arr_delay dep_delay  
##   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>  
## 1     9 370.0441    1400     227      11        2  
## 2    16 374.2731    1416     227      20        4  
## 3    31 408.3750    1089     160      33        2  
## 4   -17 516.7213    1576     183     -18       -1  
## 5   -19 394.1379     762     116     -25       -6  
## 6    16 287.6000     719     150      12       -4  
## 7    24 404.4304    1065     158      19       -5  
## 8   -11 259.2453     229      53     -14       -3  
## 9    -5 404.5714     944     140      -8       -3  
## 10   10 318.6957     733     138       8       -2
```

```
## # ... with 336,766 more rows
```

```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 x 2
```

```
##   gain    speed  
##   <dbl>   <dbl>  
## 1     9 370.0441  
## 2    16 374.2731  
## 3    31 408.3750  
## 4   -17 516.7213  
## 5   -19 394.1379  
## 6    16 287.6000  
## 7    24 404.4304  
## 8   -11 259.2453  
## 9    -5 404.5714  
## 10   10 318.6957
```

```
## # ... with 336,766 more rows
```

```
summarise(flights,  
  delay = mean(arr_delay - dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
```

```
##   delay  
##   <dbl>  
## 1 -5.659779
```

```
sample_n(flights, 20)
```

```
## # A tibble: 20 x 19
```

```
##   year month day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
## 1  2013     9   19    2105           1805        180    2217
```

```
## 2 2013 3 16 1558 1605 -7 1741
## 3 2013 6 25 2217 2030 107 2340
## 4 2013 11 16 1128 1119 9 1432
## 5 2013 12 15 1013 1019 -6 1331
## 6 2013 8 10 558 602 -4 737
## 7 2013 4 7 1435 1200 155 1646
## 8 2013 7 23 NA 1500 NA NA
## 9 2013 6 21 1733 1725 8 2024
## 10 2013 1 19 759 759 0 904
## 11 2013 7 6 853 853 0 1125
## 12 2013 12 15 926 900 26 1232
## 13 2013 10 10 559 600 -1 817
## 14 2013 2 20 1957 2000 -3 2111
## 15 2013 2 23 915 920 -5 1420
## 16 2013 11 9 2354 2359 -5 417
## 17 2013 1 5 1746 1745 1 2041
## 18 2013 7 14 1414 1425 -11 1545
## 19 2013 10 19 825 825 0 1136
## 20 2013 8 18 1529 1514 15 1638
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
sample_frac(flights, 0.001)
```

```
## # A tibble: 337 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013    11    21    1700           1700           0     1821
## 2 2013    12     2    2130           2100          30     2241
## 3 2013    12    20    2257           2141          76      152
## 4 2013     9    30    1146           1154          -8     1251
## 5 2013     4     4    2102           2059           3       17
## 6 2013    12    18    1308           930         218     1836
## 7 2013     6    29    1909           1826          43     2139
## 8 2013     5    21    1632           1529          63     1719
## 9 2013     6     4    1827           1830          -3     2200
## 10 2013     1     7    1455           1455           0     1623
## # ... with 327 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## Grouped operations

At the most basic level, you can only alter a data frame in five useful ways:

1. you can reorder the rows (`arrange()`),
2. pick observations (`filter()`) and
3. variables of interest (`select()`),
4. add new variables that are functions of existing variables (`mutate()`) or
5. collapse many values to a summary (`summarise()`). The remainder of the language comes from applying the five functions to different types of data, like to grouped data, as described next.